

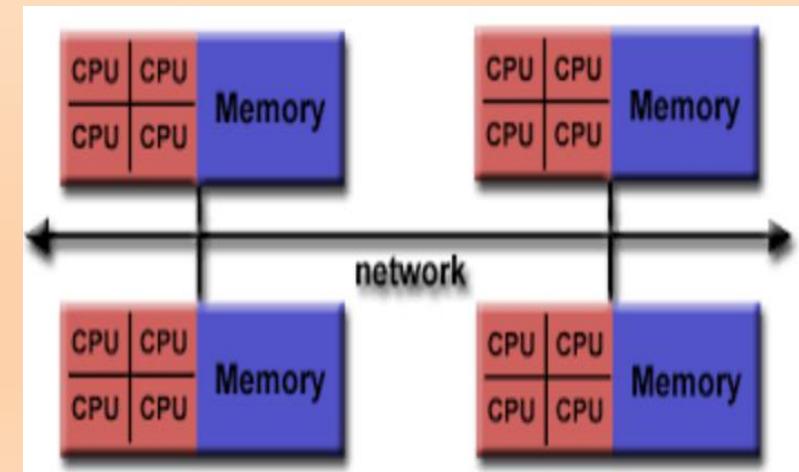
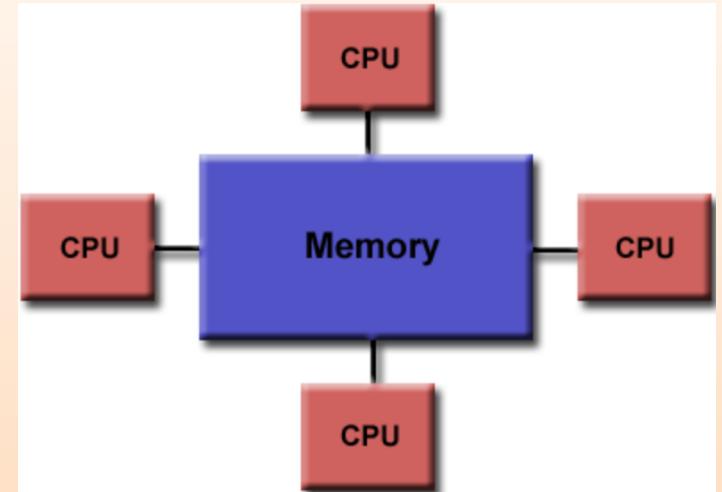


Pencil Code: Parallelization

Multi-process and multi-thread

Why does PC need multiprocess parallelization?

- medium-size setups, say 1024^3 grid points, 8 variables, double precision
 - > 128+ Gbyte main memory < on-node memory in big clusters
 - > **shared memory, one-process approach**
(but time to solution with 64 ... 128 cores?)
- large-size setups, say 4096^3 grid points, ...
 - > 8+ TByte main memory >> on-node memory in big clusters
 - > **distributed memory approach** necessary -> **multi-process** implementation
 - naturally:
 - SPMD paradigm** (Single Program Multiple Data)
 - = multiple copies of same program work on equally sized problem parts
 - PDE: update of variables in a point needs information from **local surroundings**
 - > interprocess communication needed -> **MPI**



What is Message Passing Interface (MPI)?

- standardised library - different implementations ([OpenMPI](#), [MPICH](#))
- allows **different UNIX processes** to communicate
- -> in general: Multiple Program Multiple Data paradigm ([MPMD](#)),
in particular **SPMD**: same executable runs in multiple instances
- messages = **tagged data parcels** exchanged within **communicators**



<https://www-lb.open-mpi.org/software/ompi/v5.0/>

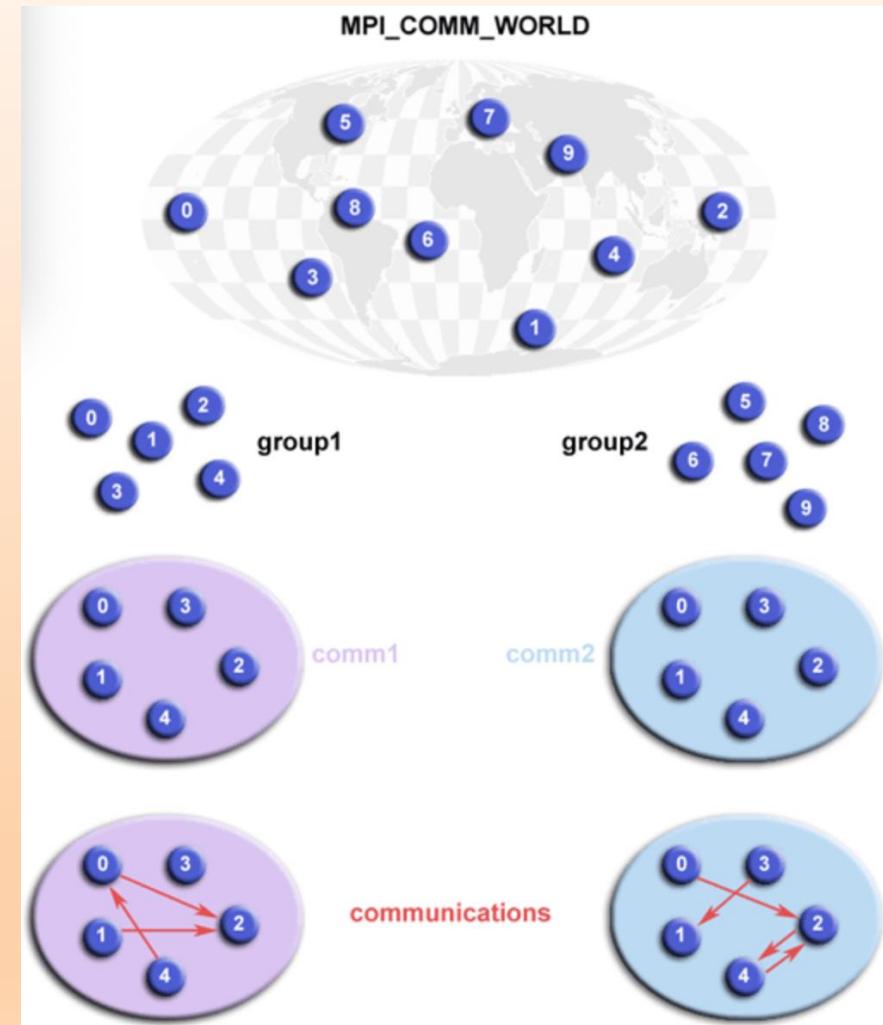


<https://www.mpich.org/documentation/guides/>



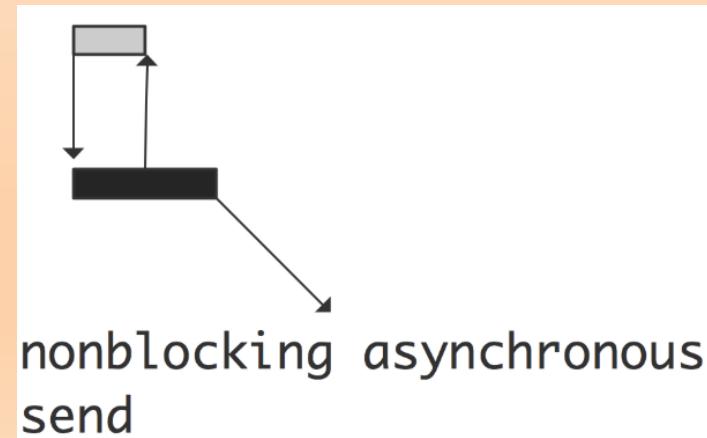
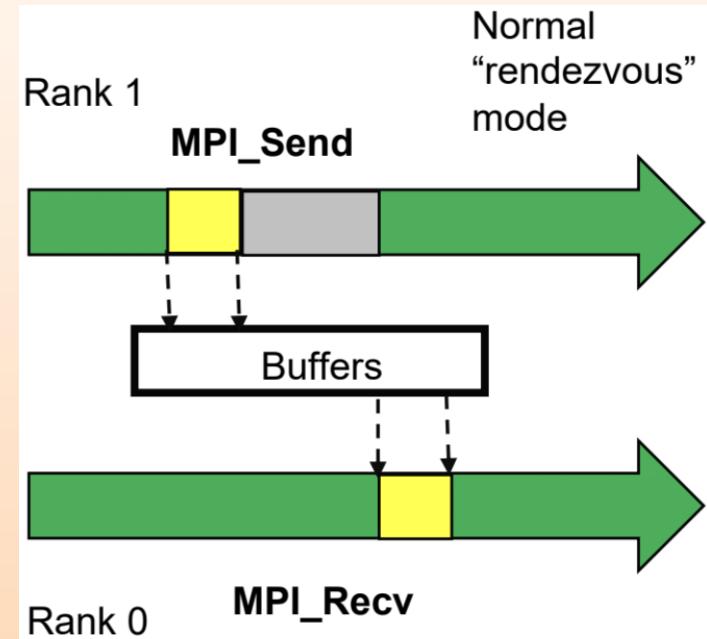
What is a communicator?

- **group** of processes with **context** (tag, topology, attributes)
- **exclusive**: messages sent within one comm. cannot be received in another comm. (exception: intercommunicators)
- **default**: **MPI_COMM_WORLD**
 - = all processes launched together by
mpirun, mpiexec, srun etc.



Point-to-point vs. Collective communications

- Point-to-point (P2P): one sending proc -- one receiving proc
(but simultaneous send and receive with one call `MPI_SENDRECV`)
- collective: **all** procs of a comm. must execute the call,
typically **reductions** like sum or max/min
- **blocking** vs. **non-blocking** communication calls:
 - blocking: return only if a "success criterion" is fulfilled (`MPI_[B|R|S]SEND`)
 - non-blocking: return immediately (`MPI_I[B|R|S]SEND`)
-> key for **concurrency of communication and computation** (or other)

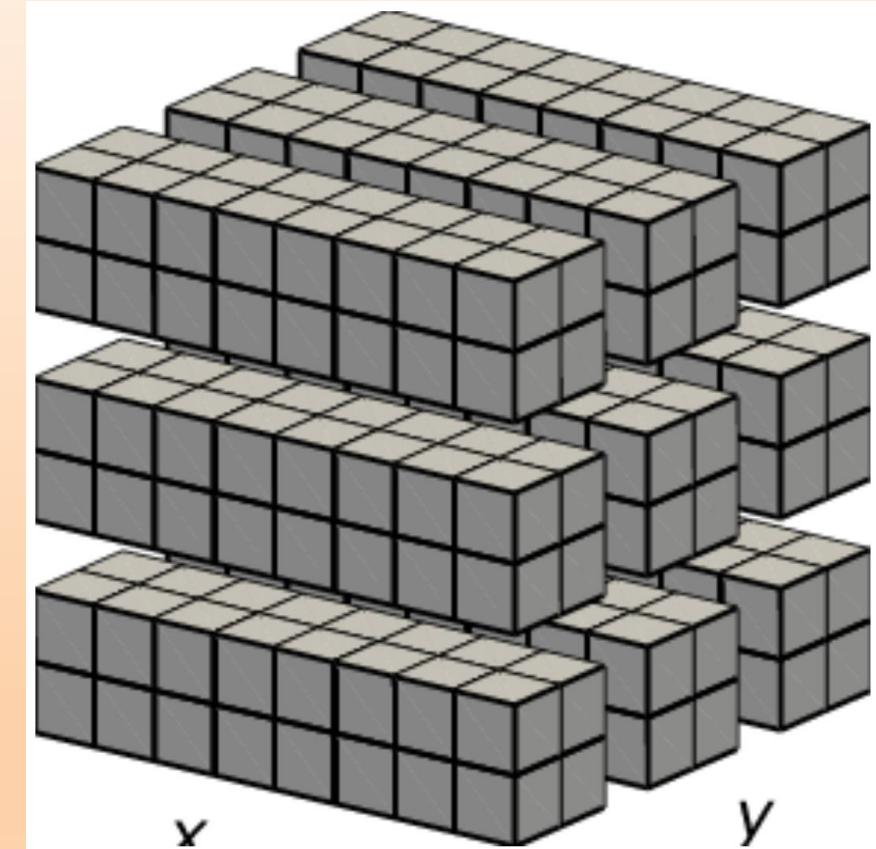


Implementation in Pencil Code

- processes arranged as a Cartesian "process grid",
defined by the user through `nproc[xyz]` in `cparam.local`
-> each process has "coordinates" `iproc[xyz]` in this grid
(Cartesian topology could be enforced by MPI tools, but is not at the moment)
- computational grid is divided in `ncpus` equally sized
cuboid **subdomains** according to `nproc[xyz]`
subdomain size:
$$(\text{nxgrid}/\text{nproc}_x, \text{nygrid}/\text{nproc}_y,$$

$$\text{nzgrid}/\text{nproc}_z)$$
- each process holds a subdomain

Example: 8x6x6 grid



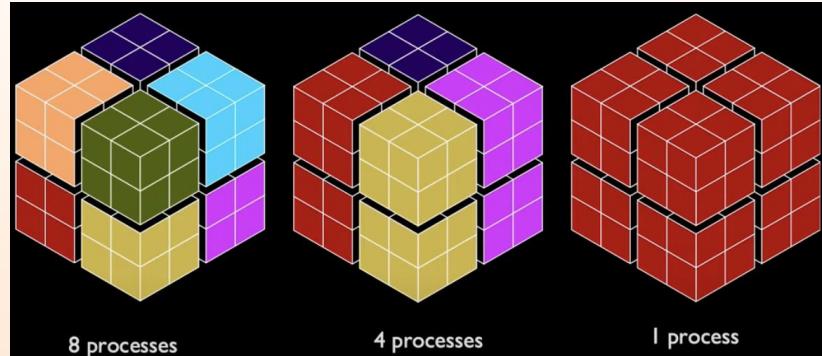
`nprocx=1, nprocy=3, nprocz=3`

Implementation in Pencil Code

- task:
mapping process numbers (“ranks”, in PC: `iproc`) to subdomains
processes linearly numbered: 0, 1, 2, ... total number of procs-1

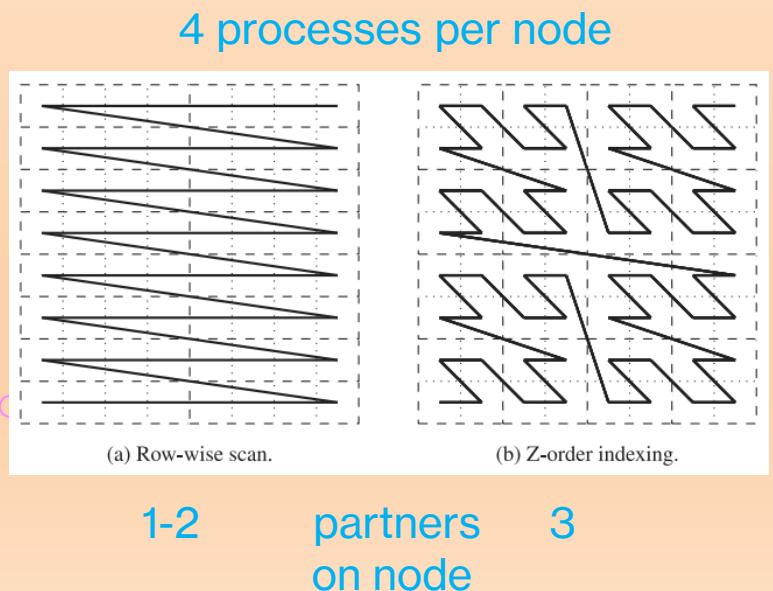
how to map `rank -> (iprocx, iprocy, iprocz)` ?

- optimal: union of the subdomains of all processes on
a compute node as close to a cube as possible
-> maximized on-node communications



- at the moment naive linear map with
`iprocx` running fastest, `iprocz` slowest:
$$\text{iproc} = \text{iprocx} + \text{iprocy} * \text{nprocx} + \text{iprocz} * \text{nprocx} * \text{nprocy}$$

- better: Morton-curve numbering



Major communication task: from stencils

- stencil structure of finite difference formulae:

- for most differential operators

3-dimensional von-Neumann stencil with radius r=3

(default, other options: 1, 2, 4, 5)

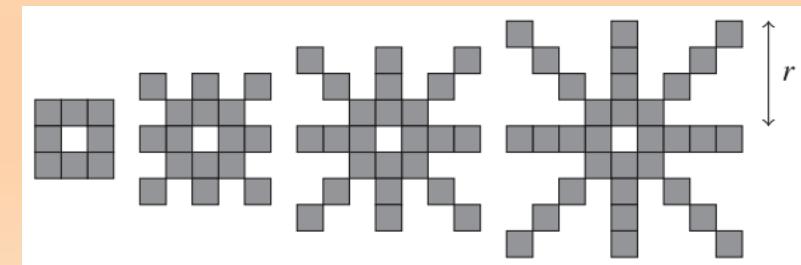
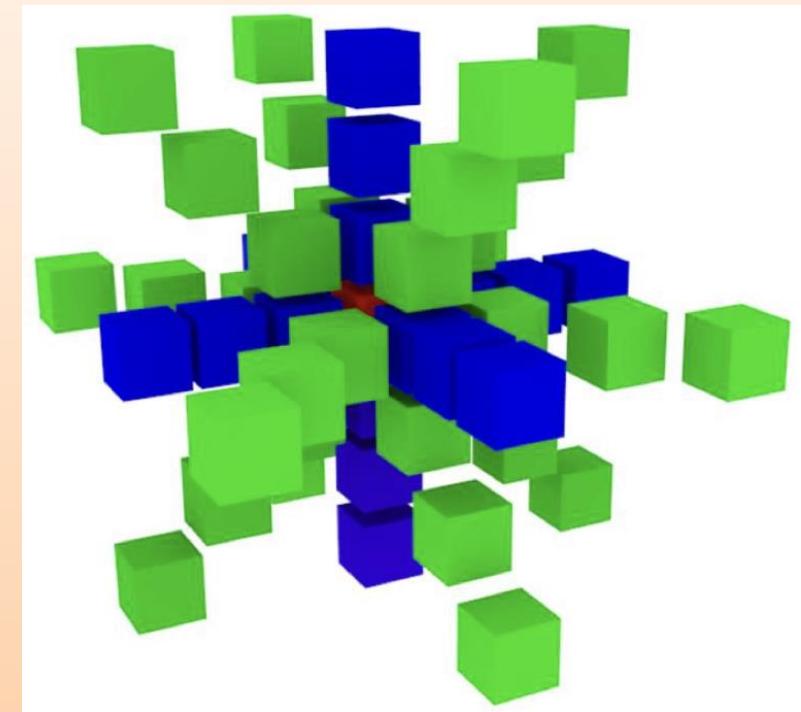
for mixed derivatives $\frac{\partial^2}{\partial x_i \partial x_j}, i \neq j$:

2-dimensional Moore stencils in xy , xz and yz planes,

but simplified to

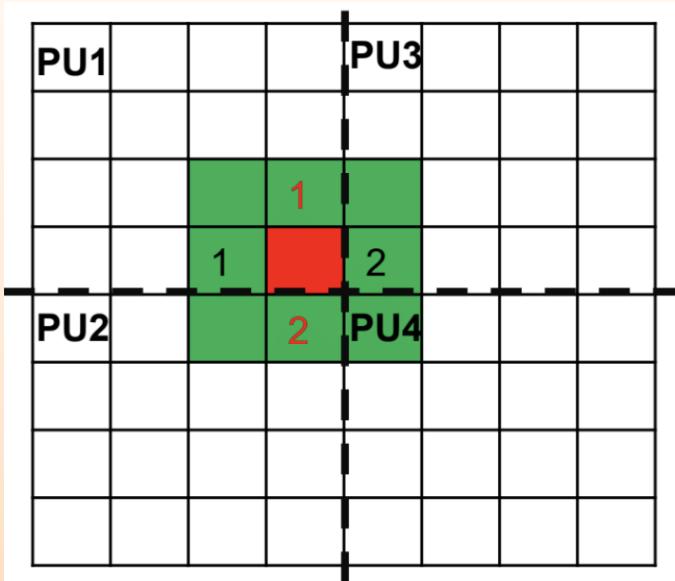
2-dimensional von-Neumann stencils w. 45° rotation

default 55-point 3D stencil in PC



2D PC-stencil r =1,2,3,4

Moore r=1 stencil



Major communication task: halo update

- stencil operations
 - > per-process grid subdomains need a **halo** ("ghost zones")
outside subdomain = **inside** subdomains of other processes
 ("inner halo")
- dictates "halo update" or "ghost zones update" through **interprocess communication**

for each variable updated by stencil operations

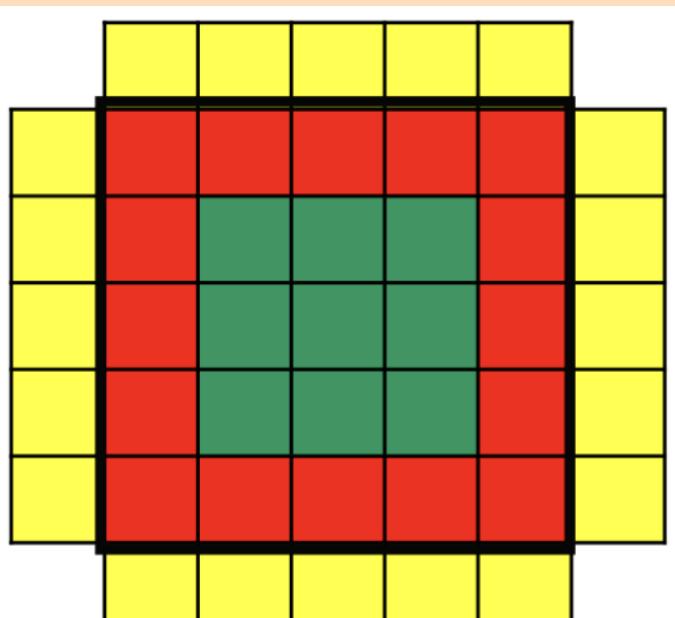
- all **dependent** variables of the PDEs
- all **auxiliary** variables subject to stencil operation (smoothing, maximum ...)
- > distinction between communicated and non-communicated auxiliaries

MAUX CONTRIBUTION

--

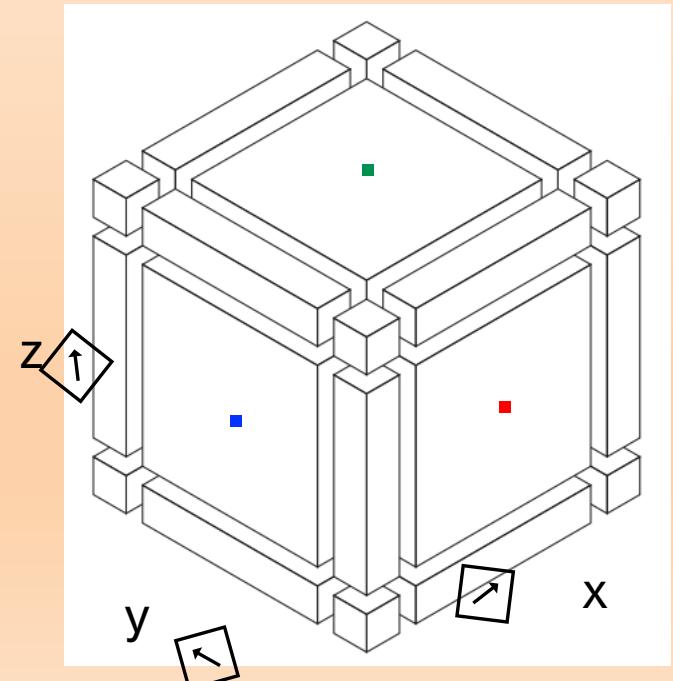
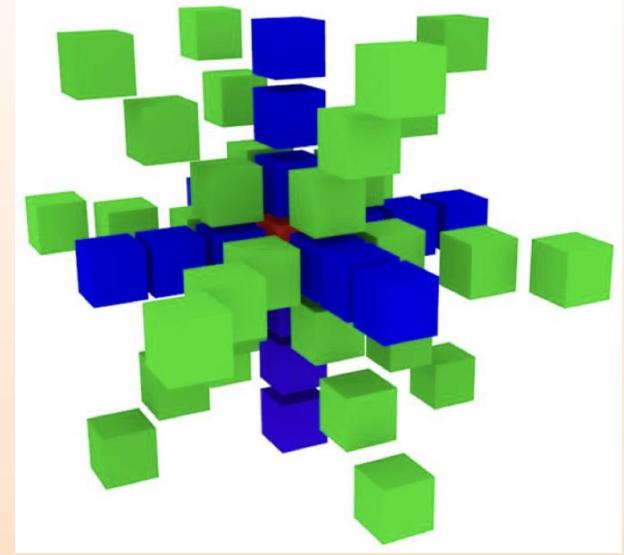
COMMUNICATED AUXILIARY

von Neumann r=1 stencil



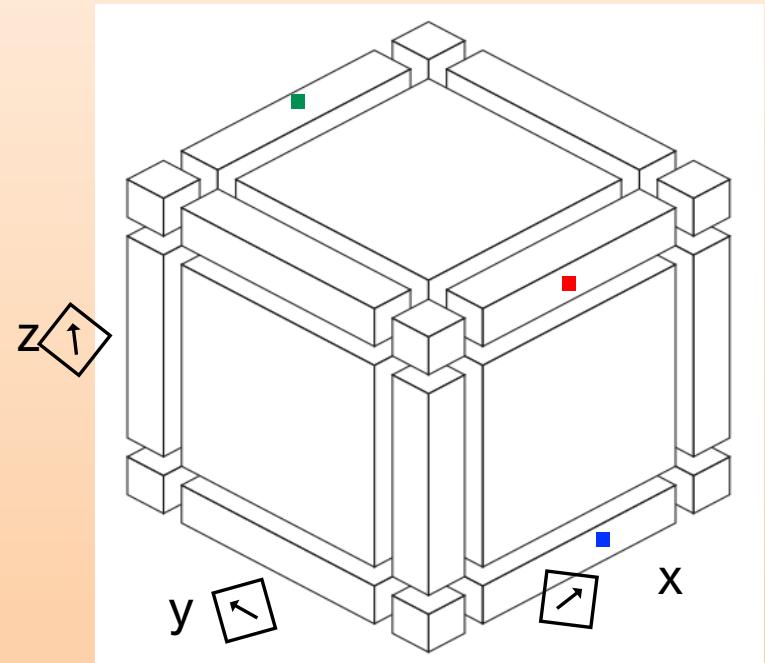
Halo update

- correspondingly: communication with
 - cross-face
 - cross-edge
- process neighbours, not with cross-vertex neighbours
- -> 6 cross-face process neighbours, communication conveys
 - left and right yz -plates in positive and negative x -direction
 - front and back xz -plates in positive and negative y -direction
 - top and bottom xy -plates in positive and negative z -direction
- -> 12 cross-edge process neighbours, theoretically



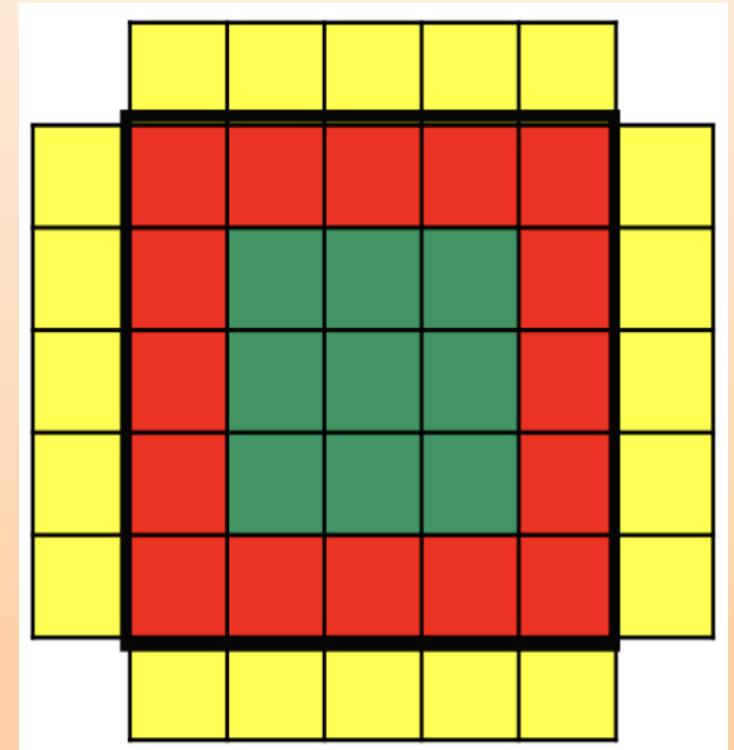
Halo update

- but only 4 cross-edge communications = those in yz planes
 - 4 x-beams in lower-lower $(-y, -z)$
 - lower-upper $(+y, -z)$
 - upper-upper $(+y, +z)$
 - upper-lower $(-y, +z)$
 - directions (yz plane)
- remaining beams (y and z):
 - cross-edge neighbour = cross-face neighbour of cross-face neighbour
-> finish x -communication first,
 - let xy and xz plates also comprise ghost zone in y and z
 - no communication across faces/edges of the global domain,
unless they belong to a periodic direction
 - instead: application of physical boundary conditions



SPMD program flow with halo communication

- in each integration substep
 - fill communication buffers with data from "inner halo"
 - initialise point-to-point non-blocking communication
 - = `call MPI_ISEND` halo-section-specific tag for all neighbours
 - call corresponding `MPI_IRecv`
 - update **inner part** of subdomain (independent of **halo**)
 - finalize point-to-point non-blocking communication
 - = `call MPI_WAIT` for each `MPI_ISEND/MPI_IRecv`
 - after completion, update **outer part of subdomain = inner halo**
 - synchronise processes (`MPI_BARRIER`) and repeat



Further MPI-parallelized tasks

- reductions
 - averages/integrals over volumes, planes or lines (=0D to 2D averages)
and global extrema for **diagnostics**:
low cadence -> cheap
 - 1D or 2D averages for special integro-differential equations
(\nearrow test methods):
every substep -> expensive
 - 0D averages for **conservation-preserving measures**:
remove mass/flow/momentum drift:
adjustable cadence -> medium exp.
 - for separately studying large- and small-scale instabilities:
~
 - for special SGS models acting only on fluctuations:
every substep -> expensive
 - implemented by blocking collective MPI calls
- spectra as diagnostics or for Poisson solving/BCs: **parallelized FFT**
cheap -> expensive
- parallel write into **monolithic snapshot file**

MPI in Pencil Code: mpicomm.f90

- encapsulates all calls into MPI library
 - > for non-parallelised runs
 - `nompicomm.f90` is included in the build instead
 - implements all preparatory work:
 - initialising/finalizing MPI
 - creating communicators
 - establishing neighbours
 - establishing "boundary process" flags
 - `l[first|last]_proc_[xyz]`
 - allocating communication buffers and tags

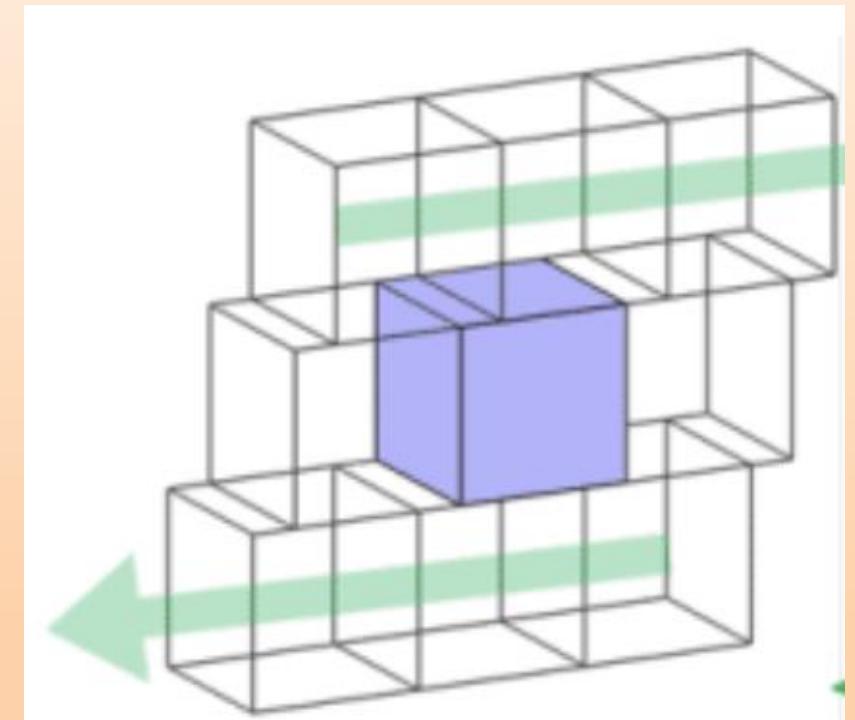
```
public :: mpirecv_logical,  
         mpirecv_real,  
         mpirecv_int,  
         mpirecv_char !,  
         mpirecv_cmplx  
         mpisend_logical,  
         mpisend_real,  
         mpisend_int,  
         mpisend_char,  
         mpisendrecv_real,  
         mpisendrecv_int,  
         mpireduce_sum_int,  
                     mpireduce_sum  
         mpireduce_sum_double  
                     mpireduce_max,  
         mpireduce_max_int,  
                     mpireduce_min  
         mpiallreduce_max,  
         mpiallreduce_min  
         mpiallreduce_sum,
```

Pencil Code communicators

- `MPI_COMM_WORLD` (default MPI-provided)
 - in PC only used at initialisation and for MPMD
- `MPI_COMM_PENCIL`
 - by default duplicate of `MPI_COMM_WORLD`, but can differ in MPMD mode
- `MPI_COMM_GRID`
 - by default duplicate of `MPI_COMM_PENCIL`,
 - but split in two if PC holds two congruent grids ([Yin-Yang grid](#))
- `MPI_[XY|XZ|YZ]PLANE` as many as there are `XY|XZ|YZ` planes in processor grid
- `MPI_[XYZ]BEAM` as many as there are `X|Y|Z` beams in the processor grid

Non-standard communication patterns: shear

- for “sliding periodic” boundary conditions (shearing-box):
up to five neighbours at x-boundaries along y-direction

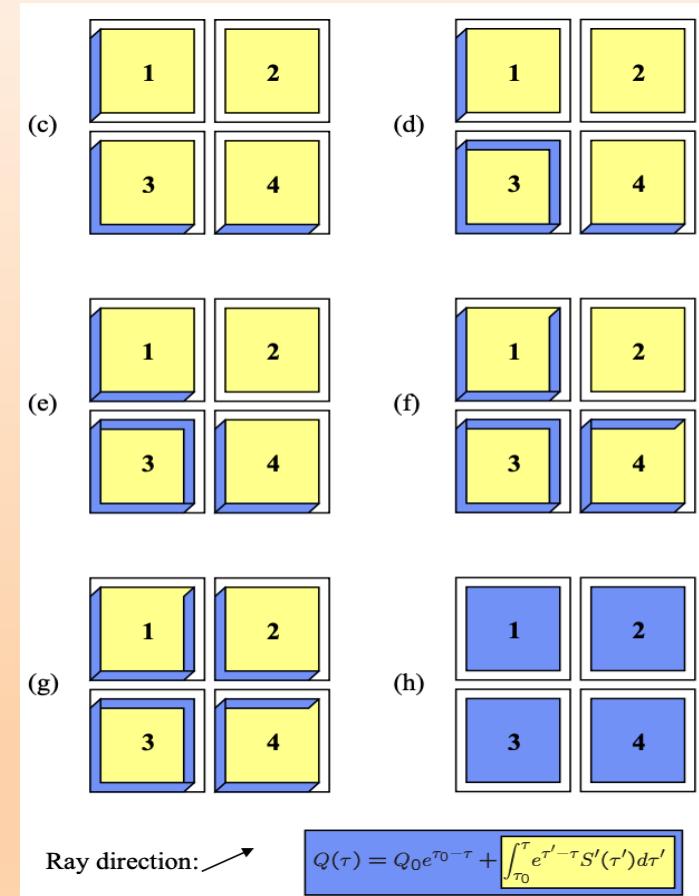


Non-standard communication patterns: radiative transport

- rays traverse whole domain, possibly several times for periodic BCs
- contributions to energy equation: **line integrals** along ray

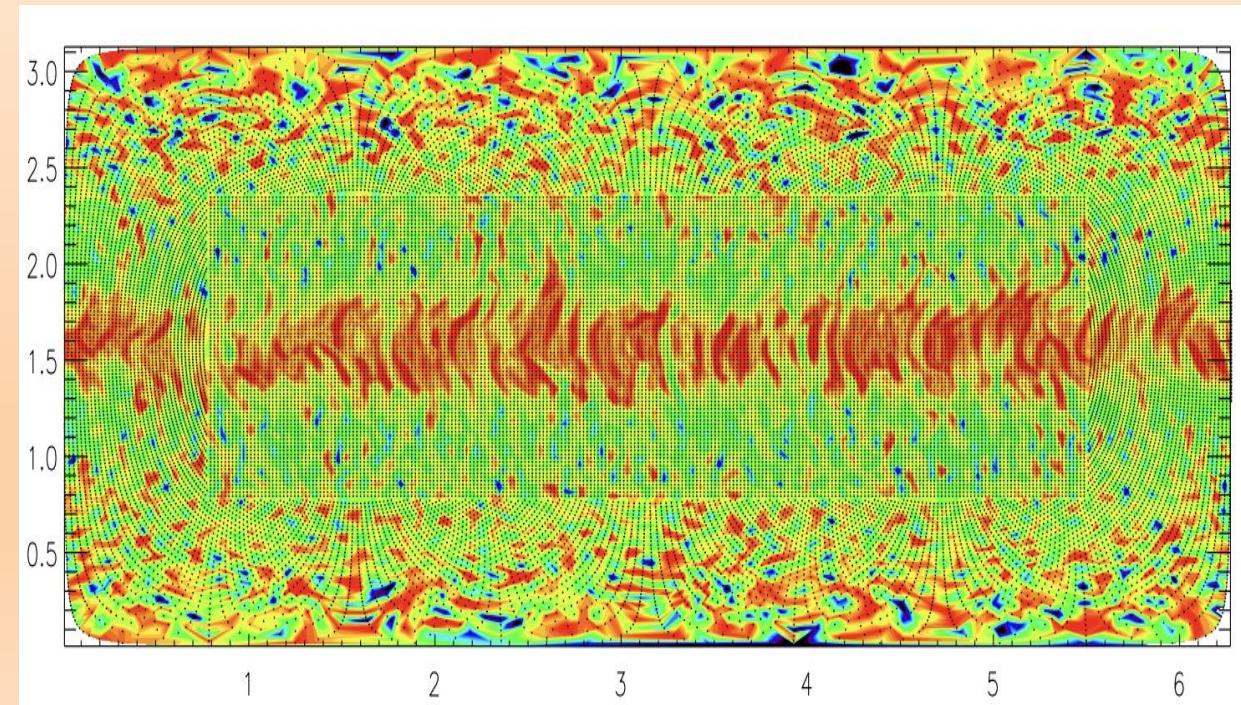
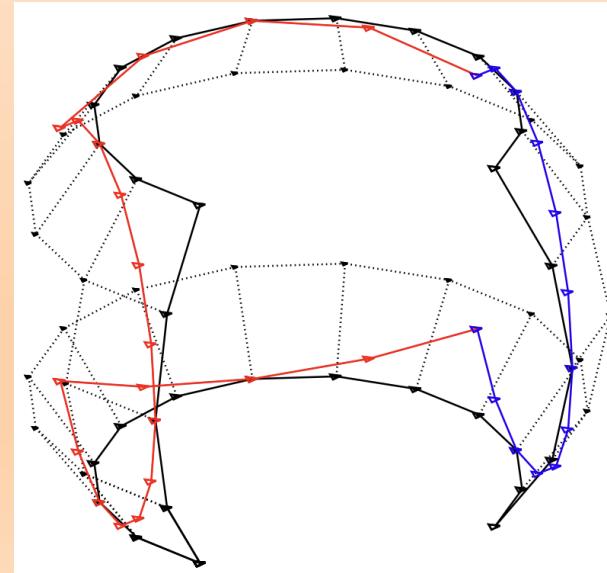
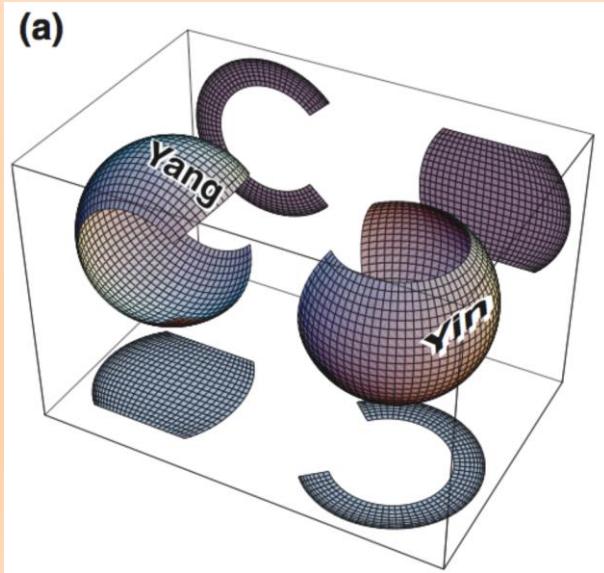
$$I(\tau) = I(\tau_0)e^{-(\tau-\tau_0)} + \int_{\tau_0}^{\tau} e^{-(\tau-\tau')} S(\tau') d\tau' .$$

-> reductions needed



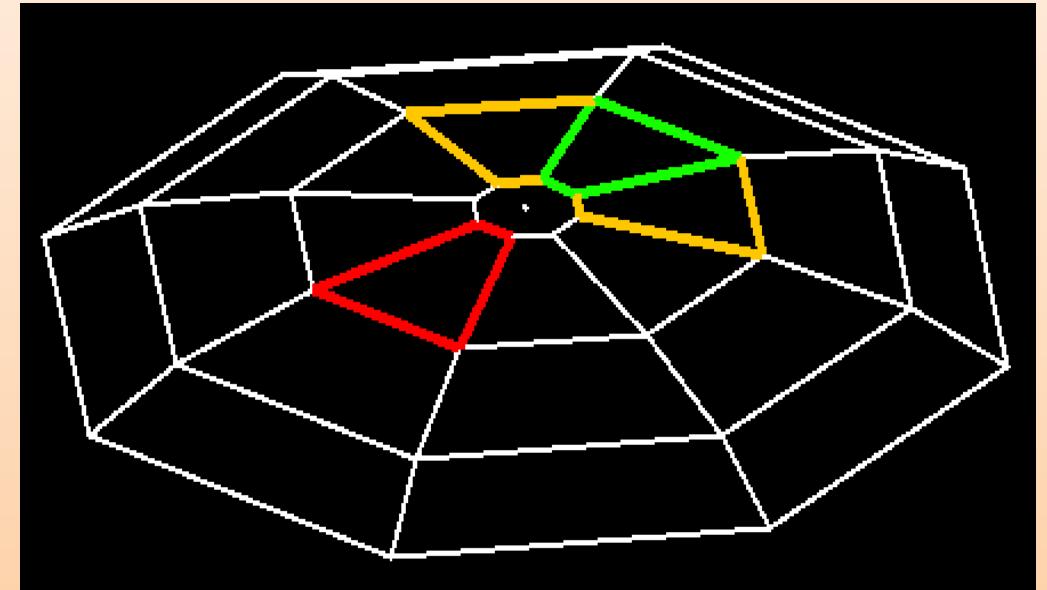
Non-standard communication patterns: Yin-Yang grid

- two congruent grids covering full sphere (communicator `MPI_COMM_GRID`)
- inside each grid: standard communication pattern
at edges: defined by grid overlap
- yet experimental



Non-standard communication patterns: cross-pole

- 3D setups in spherical coordinates (r, ϑ, φ) :
axis singularity because of metric coefficient $1 / r \sin \vartheta$
- can be avoided by $d\vartheta/2$ distance from pole
-> a ϑ boundary?
better not! artificial!
- alternative: consider “periodicity” **across the pole**
-> all “ y = cross-pole neighbors”
have same ipy proc coordinate, but different ipz
- tb combined with **grid coarsening** near poles – yet experimental

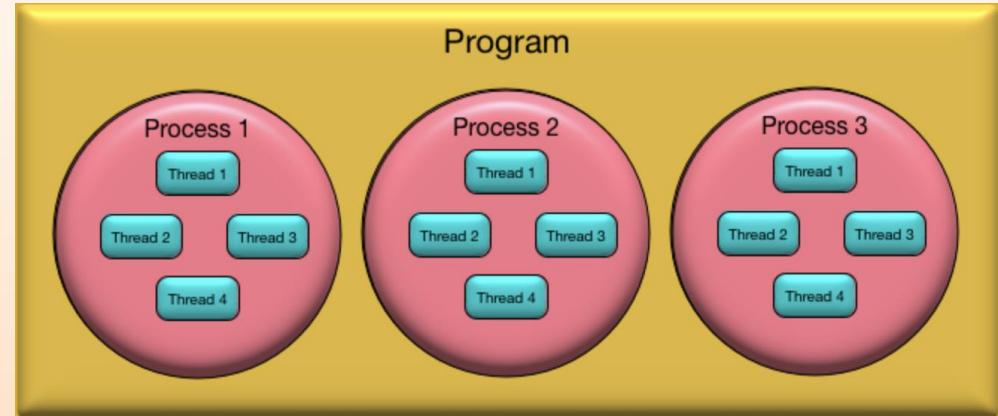


Non-standard communication patterns: MPMD

- interface to communicate with “foreign” grid-based code with independent processor layout
 - gathers information about foreign grid and layout
 - establishes communication pattern with foreign code
 - receives and interpolates data in full domain
 - yet experimental

Multithreading

- enables direct utilization of shared memory on node
(MPI silently takes advantage of it, too)
- instead of multiprocess
 - > **parallel threads of a single process** on a group of CPU cores
- meaningful in connection with GPU acceleration:
 - for example: **8 GPUs, 64 CPU** cores per node
 - > **8 MPI processes with 8 threads each**, using all **64 cores**
 - accelerates **diagnostics calculation** on CPU
 - allows **concurrence** of diagnostics calculation and output with integration on GPUs



Multithreading

- implementation: **nested parallelism**
 - first parallel region with 2 threads: master <-> “diagmaster”
master administers program flow, especially kernel launch on GPU
diagmaster ~ diagnostic calcs and output of diagnostics and snapshots
spawns new parallel region with all remaining available threads
 - encoded in Fortran by **OpenMP** directives (`!$omp ...`) and guards (`!$...`), activated by proper compilation flags
 - runtime specifications, e.g. in **SLURM**:

```
#SBATCH --cpus-per-task=8  
  
...  
  
export OMP_NUM_THREADS=${SLURM_CPUS_PER_TASK}  
export OMP_MAX_ACTIVE_LEVELS=2  
export OMP_PROC_BIND=close,spread  
export OMP_WAIT_POLICY=PASSIVE  
  
pc_run
```

