


INTRODUCTION TO PC-A

“Pencil Code with Astaroth inside”

BASIC GPU ARCHITECTURE





sequential execution

optimising


parallel throughput

10x higher peak throughput
10x higher DRAM access speed


GPU HARDWARE ESSENTIALS



GPU MEMORY HIERARCHY



GPU EXAMPLE: PASCAL GPU




CPU-GPU WORKCYCLE

in every substep:

load program (kernel) and data (for PC: a few)

wait or act concurrently
during kernel execution



store results (for PC: a few)

PC-A: MINDSET

- two different codes in two different repositories
(although sharing numerical methods)
- two different computers, although GPU dependent on CPU
→ everything needed for advancing the PDE variables
 - must be in GPU memory (parameters, initial conditions)
 - in kernel code (recipes for right hand sides, boundary conditions, etc.)
- f-array is updated on CPU only when needed, df-array does not exist

PC-A tries to use all resources on a node with maximum concurrency!

HOW TO GET THE CODES

- both codes together with a **fresh** pull:

```
git clone --recurse-submodules https://<username>@github.com/pencil-code/pencil-code.git.
```

(read-only)

or

```
git clone --recurse-submodules https://<username>@pencil-code.org/git/ pencil-code  
source sourceme.sh  
cd $PENCIL_HOME/src/astaroth/submodule  
git checkout develop
```

- or add Astaroth and Fortran_parser to an **existing** PC installation:

```
cd $PENCIL_HOME  
git checkout master (if necessary)  
git submodule update --init --recursive  
cd src/astaroth/submodule  
git checkout develop
```

BRANCHES MATTER!

HOW TO BUILD PC-A

- on clusters, load appropriate CUDA (NVIDIA) or HIP (AMD) and MPI modules, perhaps `cmake`
- set in `Makefile.local`

```
GPU           = gpu_astaroth
MULTITHREADING = openmp
RUNTIME_COMPILATION = on
TRANSPILATION      = on
GENERATE_DSL_CODE = on
```

- if you haven't used PC-A before in work directory, execute `pc_setupsr`

-> new symbolic links and directories in

`src/astaroth` interface routines and PC specific DSL code

`src/astaroth/submodule` Astaroth library and DSL compilation framework

- build Pencil Code as usual:

- `make`: set `MODULE_ [PRE|IN|SUF] FIX` environment variables as given in config file
- `pc_build`: use proper config file (flag `-s | --serial` for troubleshooting)
- `gfortran`: set `FFLAGS+=-mcmodel=[medium|large]` for large grids
- `CRAY`: set `FFLAGS+=-h pic -WI,--no-relax`

HOW TO BUILD PC-A

- build process creates
 - DSL and interface code according to the setup
 - Astaroth libraries and interface library
 - for libraries: separate `Makefile` in `src/astaroth`
-> libraries can be built **separately** (perhaps after "`make clean`", answer "y")
- most importantly: **kernels and their grouping** in `src/astaroth/DSL/local`:
`mhd solver.ac`
- which includes
`steps_two.h`
`boundconds.h`: boundary conditions (mandatory "step")

HOW TO RUN PC-A

Example SLURM batch script:

```
#SBATCH --nodes=2          # Total number of nodes
#SBATCH --ntasks-per-node=8
#SBATCH --cpus-per-task=7   # multithreading

CRAY (LUMI, Dardel, Frontier):
#SBATCH --gpus-per-node=8   # Allocate one gpu per MPI rank
```

CSC machines:

```
#SBATCH --gres=gpu:v100:4      ~
```

```
source src/.moduleinfo

export LD_LIBRARY_PATH=${CRAY_LD_LIBRARY_PATH}:$LD_LIBRARY_PATH
export OMP_NUM_THREADS==${SLURM_CPUS_PER_TASK}
export OMP_MAX_ACTIVE_LEVELS=2
export OMP_PROC_BIND=close,spread
export OMP_WAIT_POLICY=PASSIVE
./start.csh
export MPICH_GPU_SUPPORT_ENABLED=1 # MPICH only
./run.csh
```

ALTERNATIVES: OFFLOADING

- OpenACC:

```
#pragma acc data copyin(a[0:n], b[0:n]) copyout(c[0:n])
// defines the data that needs to be moved to the device
{
    #pragma acc parallel loop.                                //
offloads the loop for parallel execution
    for (int i = 0; i < n; ++i) {c[i] = a[i] + b[i];}
} // End of data region. The data is copied back from the device.
```

- OpenMP:

```
#pragma omp target "map(tofrom:y[0:n])" //other variables implicitly
copied to device
#pragma omp teams distribute parallel for
for (int i = 0; i < n; i++) { y[i] = a * x[i] + y[i]; }
```

- Others: Kokkos (library calls).

WHAT IS ASTAROTH?

- library for creating efficient GPU kernels for stencil operations
 - analyzing task dependencies
-> launch kernels concurrently using CUDA streams
 - holding intermediate data in caches for reuse
 - autotuning for optimal thread block sizes
- Domain-Specific Language (DSL) and compiler
 - facilitates writing of kernels
 - detects communication needs
 - establishes task graph

WHAT IS DOMAIN-SPECIFIC LANGUAGE?

- C-like, with some Python-like features (partly declaration-free)
- Example: rhs of continuity equation

Field LNRHO

```
#define RHO LNRHO
Field3 UU
...
dlnrho_dt() {
    glnrho = gradient(LNRHO)           // grad(rho) or
grad(lnrho)
    rhs = - dot(UU, glnrho)
    if (ldensity_nolog) {return    rhs - RHO*divergence(UU) }
    else
        {return    rhs - divergence(UU) }
```

WHAT IS DOMAIN-SPECIFIC LANGUAGE?

- operations to be understood **pointwise on grid** (one CUDA thread per grid point)
- essential types:
 - `Field, Field3` array on grid subjectable to stencil operations
 - `Field chemistry[n_species]` many-species array
 - `Kernel()` transforms into CUDA Kernel
 - only way to call DSL code
 - `Stencil()` compact stencil definition
- qualifiers:
 - `dconst` on device constant memory (fast)
 - `run_const` constant during runtime (optimized away)
 - `gmem` on device global memory (for array-type parameters)

PARAMETER TRANSFER TO GPU

- parameters of physics modules are **pushed to GPU**

by `subroutine pushpars2c(p_par)` near end of each physics module:

```
call copy_addr(<parameter>, p_par(<running index>))
```

for a new parameter: add a `call copy_addr(...)`, increase `n_pars` if needed

`integer parameters:` add ! int

`logical ~` add ! bool

`real array ~` add ! (<dim>) [(<dim>) ...] at line end

- note: parameter manipulations, doable in module initialization,

not to be coded in DSL

-> push **derived parameters** (typically logicals)

- parameters from `src/cparam.h`, `src/cdata.h`: all available

CONTROL FLOW DESIGN

- in general: all operations modifying **f-array** should happen on GPU
-> branches `if (lgpu) then ...` at all relevant places in PC sources, e.g.:

```
call before_boundary_gpu
call rhs_gpu
```
- called routine finally calls `acGridExecuteTaskGraph(<name of compute step>,...`
- **task graph**: organizes execution of kernels in maximally concurrent way,
inferred from task dependencies, derived from “**compute step**”
- communication not explicitly specified:
can be inferred from the modification of the “output buffer”
- specifically, **concurrency** of communication and computation is **maximized**

COMPUTE STEPS

- **compute step**: collection of **kernels**, to be executed **without interference of CPU**, presently:
 - AC_rhs
 - AC_calculate_timestep
 - AC_before_boundary_steps
 - AC_before_boundary_steps_including_halos
 - AC_after_timestep
 - AC_gravitational_waves_solve_and_stress
 - AC_calc_selfgravity_rhs
 - AC_sor_step
 - need to be extended by hand
-
- The diagram illustrates the grouping of compute steps. It features two pink curly braces. The first brace groups the 'AC_calc_selfgravity_rhs' and 'AC_sor_step' items under the label 'for self-gravity'. The second brace groups the 'get_source_function_and_opacity', 'Qintrinsic_steps', and 'Qextrinsic_steps' items under the label 'for radiation'.

TRANSPILATION

- handwritten DSL code: error prone, needs human interference
for
keeping in sync with Fortran
- alternative – **transpilation**:
transforms all code employed in `rhs_cpu` into **single DSL kernel**
 - all used module variables become global, in const or global GPU memory
 - **Pencil Case** dismantled into scalar variables,
all “penciled” variables and operations -> **pointwise**
 - built-in functions of Astaroth (`stdlib`) used instead `sub.f90`, etc.
 - **f-array** -> individual **Field** objects, df-array: local variables
- can be part of PC-A build process

TRANSPILATION

- other transpilable cases, not part of PC-A build process:
 - boundary conditions
 - `before_boundary`, `after_boundary` stuff
- workflow: customize `parse.py` (`$PENCIL_HOME/fortran_parser`)
 - transpile by hand (`python transpile.py`)
 - set preprocessor guards (e.g., `#if LHYDRO`)
 - add `reduction calls` by hand if needed
- new fields: handwritten addenda to
 - `fieldecs.h`, `df_declares.h`, `handwritten_end.h`
- commit transpiled code to `src/astaroth/DSL`

OPTIMIZATION: RUNTIME-COMPILATION

- for optimization, Astaroth needs to know which **stencil operations** and **reductions** are needed
- PDE rhss contain many **conditionals**, depending on **runtime data**
-> optimization inhibited
- solution: turning **run-time constants** into **compile-time** constants
-> re-compilation of DSL code at run-time after reading of input

RUNTIME-COMPILATION

- elimination of all conditionals
-> totality of all stencil calculations can exactly be inferred
- for boolean/integer variables **changing within the time loop**:
kernel variants are produced
- unneeded Fields and arrays not allocated, unneeded compute steps eliminated
- uncommunicated fields detected -> only one buffer in GPU global memory

LIMITATIONS

- **uncovered** modules: see exclusion list in `gpu_astaroth.f90`,
`initialize_GPU`
 - **incomplete** before/after boundary stuff
 - **missing** diagnostics (when not derived from pencils)
 - only **one-processor FFT** on GPU
 - **limited testing** (see <https://norlx51.nordita.org/tests/GPU/>)
 - build process & runtime-compilation are **divas**
- > always test against CPU version !!!



but promise: speedup = 10 ... 30 !

TROUBLESHOOTING

- enforce re-creation of interface code by

```
rm src/astaroth/PC_moduleflags.h (tb improved)
pc_build ...
```
- and libraries

```
cd src/Astaroth
make clean
make
```
- obey DSL syntax meticulously - consult:
<https://bitbucket.org/jpekkila/astaroth/src/develop/acc-runtime/README.md>
- compare with samples in GPU autotest
- consult touko.puro@aalto.fi or matthias.rheinhardt@aalto.fi



OUTLOOK

- in preparation:
 - run-time compilation -> all variables constant during time-loop,
esp. logical variables, are replaced by their values from `start.in/run.in`
performance!

full transpilation of the rhss to DSL -> manual DSL coding no longer needed!

WHAT THE BUILD YIELDS

- virgin build -> void rhs functions in `equations.h` -> user intervention needed:

inspect directories in `src/astaroth/DSL/`

`density`

`entropy`

`forcing`

`hydro`

`magnetic`

`shock`

`supernova`

for useful code snippets

- indicate, which physics modules are **supported** presently
- differential operators etc. are in `src/astaroth/DSL/stdlib`

WHAT THE BUILD YIELDS

- Kernel and boundary conditions in `src/astaroth/DSL/local:`

`mhd solver.ac`, `boundconds.h`

`density`

`entropy`

`forcing`

`hydro`

`magnetic`

`shock`

`supernova`

for useful code snippets

- indicate, which physics modules are **supported** presently
- differential operators etc. are in `src/astaroth/DSL/stdlib`

CUSTOMIZE RHS

- to specify a rhs, modify e.g., function `dlnrho_dt` of `src/astaroth/DSL/local/equations.h`

```
dlnrho_dt(int step_num) { return 0. }
```

to

```
dlnrho_dt(int step_num) {
    #include "../density/continuity.h"
}
```

or

```
duu_dt(int step_num) { return real3(0.,0.,0.) }
```

to

```
duu_dt(int step_num) {
    #include "../hydro/momentum.h"
}
```

CUSTOMIZE RHS

- "physics branches" in DSL code can be selected by preprocessor statements like

```
#if LMAGNETIC  
    ...  
#endif
```

for each enabled physics module, a flag is predefined in [src/astaroth/PC_moduleflags.h](#)

- or conditionals in DSL syntax, like

```
if (ltemperature) {  
    ...  
}
```

- all switches from [src/cparam.inc](#) available
- changes to [equations.h](#) are **permanent** = not overwritten by future builds
(additional physics -> new empty rhs functions appear,
no longer needed functions do not disappear but are idle)
- [equations.h](#) is considered by [pc_newrun](#) and [cvscsi_run](#)

CUSTOMIZE RHS

- Caveat: it is advisable to check predefined DSL code in the beginning,
at least more complex functions like `denergy_dt`
- Note: to enable an additional physics module, the block in
`src/gpu_astaroth.f90` has to be released;
can require some non-standard code development
- Limitations:
 - particles/pointmasses/radiation/solid cells/self-gravity/testfields
presently not supported
 - modifications of f-array in `*before/*after_boundary` routines completely implemented
if needed every timestep
 - diagnostics, which are not only from `pencil_case/f-array` not calculated
 - not all boundary conditions "transpiled" yet (coming soon)